Supplementary Material: Choosing a Fast Stand-In for a Slow Simulation

[Author names to be supplied by the authors]

S1. Proofs

Both results are standard; we include short proofs for completeness.

Proof of Proposition 1.. Let f^{-i} minimise $\sum_{j\neq i}(y_j - f_j)^2 + \lambda N\|f\|_K^2$ and let y' equal y except $y'_i = f_i^{-i}$. For any f, the objective with data y' equals the leave-one-out objective plus $(y'_i - f_i)^2 \geq 0$, so f^{-i}, which makes the extra term zero and minimises the first, also minimises the objective with data y'. By uniqueness of the ridge solution, $f^{-i} = Hy'$. Hence $f_i^{-i} = (Hy')_i = \hat{y}_i - H_{ii}y_i + H_{ii}f_i^{-i}$, so $(1-H_{ii})f_i^{-i} = \hat{y}_i - H_{ii}y_i$ and $y_i - f_i^{-i} = (y_i - \hat{y}_i)/(1-H_{ii})$. For $K = Q\,\mathrm{diag}(w)Q^\top$ with Q orthogonal, $H = Q\,\mathrm{diag}(w/(w+\lambda N))Q^\top$ and, since $\sum_j Q_{ij}^2 = 1$, $1 - H_{ii} = \sum_j Q_{ij}^2\,\lambda N/(w_j + \lambda N)$. For the polynomial model with a thin design matrix $F = U\,\mathrm{diag}(s)V^\top$, $1 - H_{ii} = (1 - \sum_j U_{ij}^2) + \sum_j U_{ij}^2 \lambda N/(s_j^2 + \lambda N)$. A bug in an early version of the tool (Sec. S5, V9) came from computing $1 - H_{ii}$ by subtraction and clipping it. □

Proof of Proposition 2.. Let $P = V_r V_r^\top$. For $x \in \bar{\mathbf{u}} + \mathrm{span}(V_r)$, $x - \mathbf{u} = (x - \bar{\mathbf{u}} - P(\mathbf{u} - \bar{\mathbf{u}})) - (I - P)(\mathbf{u} - \bar{\mathbf{u}})$, where the first term lies in $\mathrm{span}(V_r)$ and the second in its orthogonal complement. By Pythagoras, $\|x - \mathbf{u}\|_2^2 \geq \|(I - P)(\mathbf{u} - \bar{\mathbf{u}})\|_2^2$. Dividing by $\|\mathbf{u}\|_2$ and averaging over runs gives the claim. □

S2. Pre-specification record

Hypotheses, success criteria and analysis code were committed to the project's git repository before the results they concern (a self-hosted history, not an independent registry). Table S1 lists, for each experiment, the protocol commit and the first commit that contains its results; the check is automatic (`audit/prereg_order_check.py`) and shows that no protocol was committed after its results. It shows protocol-before-results in history, not

protocol-before-first-run. Experiments H, G3 and the Stage-6 experiments T, C, S and R are labelled post hoc in motivation: they were designed in response to earlier results or to reviews (Sec. S5). For T, C and S the criteria are in commit a99c2c6; the tool was changed once afterwards, before any evaluation run, to fix a numerical bug found in smoke tests on three of the eight problems (deviation V9); experiment C ran after the results of T, including the failed prediction for H5, were known, with the pool unchanged. The real-data check R is in commit 2619b07. Three corrections to the first version of this record: (i) the analysis script for T (`analyze_s6.py`) was committed 14 minutes *after* the first T results, although the criterion and the cell definition were fixed in the addendum; for C, S and R the analysis scripts precede the results; (ii) the column "first results commit" now refers to the first commit that contains a per-job result file (an earlier version listed commits that contained only `environment.json`); (iii) several result commits (H8, H32, G3) are interim commits made by an automatic hook while runs were in progress, so the column shows git order, not run-start times, which git does not record. The search space of the tuned network was narrowed (depth 2–5, width 64–512 to depth 2–4, width 64–256) in the addendum text ten minutes after the commit of the criteria and before any run (deviation V12). The Stage 7 experiments (labels 7T, 7C, 7S, 7N, 7SP, 7E) all refer to one freeze commit, which contains the protocol addendum v3.2 with criteria and predictions, the new problems and their unit tests, the tool and the run scripts; its hash is also recorded in a separate file, and the tool had not been fitted to any of the new problems at that commit. Interim result commits were made while the runs were in progress, so the hours in the table are not run-start times.

S3. Stage 7: further problems and cost of the expensive-solver case

The eight further problems are defined in the benchmark package (one module); the unit tests check mesh convergence, scaling and symmetry identities, the Newton iteration, boundary-region selection and the location of the resonance.

S4. Earlier-stage hypotheses and verdicts

Table S4 lists the hypotheses and primary tests of the earlier stages (original 16 problems; independent replication on eight new problems; mechanism controls; longer-training and precision checks). Most confirmatory tests of

Table S1: Order of pre-specification and results commits.

Exp	Protocol commit	First results commit	Hours before	Order ok
A	bcf7f30	2d0e2d4	9.33	yes
D	bcf7f30	2abe263	10.99	yes
B	bcf7f30	72cb02c	11.07	yes
A2	72cb02c	703ad59	2.38	yes
E	72cb02c	703ad59	2.38	yes
A3	c9a7b93	39b1b7f	3.66	yes
G	c9a7b93	39b1b7f	3.66	yes
F	c9a7b93	39b1b7f	3.66	yes
H8	48a0698	84f2a21	0.08	yes
H32	48a0698	84f2a21	0.08	yes
G3	4cd411a	b37dd4e	0.05	yes
6T	a99c2c6	0e45aae	32.58	yes
6C	a99c2c6	1dc799e	53.43	yes
6S	a99c2c6	9ebaf04	56.97	yes
R	2619b07	a88ad5f	19.76	yes
7T	a096980	af431a1	0.51	yes
7C	a096980	7c4953c	1.51	yes
7S	a096980	36abd58	2.01	yes
7N	a096980	6fc0b43	2.86	yes
7SP	a096980	b22bb2e	3.46	yes
7E	a096980	5e1f2d3	5.85	yes

the first stage were not supported; the one supported at cell level (unequal tuning) rests on 5 of 16 problems.

Pilot diagnostic. A frozen rule built from the original problems was tested on eight new problems from four new families: ρ =-0.38 (exact $p = 0.18$); out-of-sample gain 0.60 with Holm-adjusted $p = 0.051$ (10^6 bootstrap draws); the threshold rule classified 6 of 24 non-tie cells correctly. The pilot statistic also correlates with the number of inputs, and all problems with $f_1 < -2$ among the 24 have three inputs; we conclude that the diagnostic is inconclusive and do not use it in the tool.

Robustness. Class labels agree between halves of the random-search draws in 83% of cells (sign agreement when both clear: 99%); 80% of reduced-basis-clear cells and 57% of network-clear cells have all three seeds beyond the tie band; class labels computed from the median instead of the mean error

Table S2: The eight further problems (`scikit-fem`); p inputs, m outputs.

Problem	p	m	Description
`helmholtz_L1`	3	256	Helmholtz (undamped), $k^2 \in [2, 20]$, which reaches the first eigenvalue $2\pi^2 \approx 19.7$ of the square; source position
`helmholtz_L2`	4	256	Helmholtz (undamped), $k^2 \in [30, 70]$, which contains the double eigenvalue $5\pi^2 \approx 49.3$; source position and width
`reacdiff_fe_L1`	4	256	nonlinear reaction-diffusion (Newton), uniform D
`reacdiff_fe_L2`	6	256	nonlinear reaction-diffusion (Newton), 2x2 blocks of D, strong reaction
`disc_elast_L1`	5	82	plane stress on a disc, 4 angular stiffness sectors + load angle
`disc_elast_L2`	9	82	plane stress on a disc, 8 angular stiffness sectors, high contrast
`advdiff_cell_L1`	3	256	advection-diffusion in a cellular flow
`advdiff_cell_L2`	4	256	advection-diffusion in a cellular flow, thin layers, source width

Table S3: Cost of the surrogate for the finely meshed plate (one solver run: 1.52 seconds). The tool time includes the cross-validated fit, the learning curve and the evaluation-only refit of all members.

	$N = 64$	$N = 256$
Test error of the selected model	3.7%	0.08%
Simulation time for the training runs (hours)	0.03	0.11
Tool time (seconds)	9	120
Predictions after which the surrogate is cheaper than the solver	70	335

agree in 95% of cells; the extrapolation test set gives a median $\bar{y}$ of -0.03 against 0.16 on the interpolation set at $N = 1024$.

S5. Deviations from the plan

Every change to the plan is recorded with its date and reason in the repository (`research_design/deviations.md` for the first stage and `deviations_v2.md` afterwards); the text is reproduced here.

First stage (original 16 problems)

D1 Run restarts (operational). First launch killed (BLAS thread oversubscription) before any job finished; the container restarted twice later and killed in-flight jobs. Completed jobs are atomic (JSON written only on completion); the runner resumes and skips them. Identical code/protocol. *Effect:* none

Table S4: Earlier-stage tests. $\bar{y}$ as in (4).

ID	Criterion	Observed	p	Verdict
H1a (v1)	median $\bar{y}$ at N=32 positive (reduced basis better)	median -0.059	0.874	not supported
H1b (v1)	$\bar{y}$ falls with N in most problems	5/16 problems	0.962	not supported
H2 (v1)	pilot features beat $\{\log_2 N, p\}$	CI [-0.07, 6.40]	0.145 (Holm)	not supported
H4a (v1)	unequal tuning moves wins toward the more tuned side	10 vs 1 discordant cells	0.006	supported (cell level)
P1 (Exp A)	Spearman($f_1, \bar{y}$)< 0, 8 new problems	$\rho = -0.38$	0.180	not rejected
P2 (Exp A)	frozen rule beats baseline out of sample	gain 0.60	0.051 (Holm)	not rejected (knife-edge)
Q1/Q2 (A2)	≥ 4 of 7 smooth candidates with $\bar{y} \geq 1$; none network-clear	7/7; 0	–	met (confounded with p)
D1	narrow-range kernel ridge within 3× of GP in $\geq 75\%$ of cells	median ratio 14.3; 13% within 3×	–	rejected
E	iso-GP vs ARD-GP and kernel ridge	ratios 1.92, 8.1	–	marginal-likelihood rule met (low p only)
G1	full-field GP within 3× of POD+GP	median ratio 0.95	–	supported
H (post hoc)	wide-range kernel ridge within 3× (B=8 or 32)	medians 6.5, 2.7	–	neither rule met (mixed)
B1/B2	4× epochs: sign kept in $\geq 85\%$; median shift < 0.5	88%; 0.24	–	met, but 15 of 48 classes change
F	float64 + L-BFGS: ≥ 5 of 7 decisive gaps remain ≥ 1 decade	5 of 7	–	met

D2 Interim analysis (disclosure). A premature analysis was run when 47 of 48 jobs existed (seeds 0,1 complete; my file count wrongly included `environment.json`). It used only complete seed batches (0,1) and was **not acted upon**; no code, hypothesis or setting changed. Output kept in `results/interim_2seed_disclosed/`. The confirmatory analysis is the 3-seed run. Verdicts of H1a, H1b, H2, H4a were identical in interim and final. *Effect:* none

D3 Seeds (protocol). Seeds 0–2 (mandatory) completed; extension seeds 3–4 **not run** (compute/time). *Effect:* declared in protocol

D4 Post-hoc exploratory analyses (exploratory). Added after seeing confirmatory results; every one is labelled *exploratory* in the manuscript and none changes a confirmatory verdict. See `results/exploratory/` . *Effect:* none

D5 Non-reproducible search sampling (bug) (disclosure + fix). Found in the fresh-clone audit: `hpo.py` seeded the random-search RNG with Python's per-process-salted `hash(method)`, so trials 1-7 of every hyperparameter search are sampled differently in every process. The experiment ran with this code. Effect: a from-scratch re-execution samples

different trial configurations than those recorded (statistically equivalent, not bit-identical). All sampled configs and their validation/test errors are stored in `results/raw/*.json`; `code/experiments/replay_check.py` refits stored best configs and reproduces 20/20 recorded test errors exactly (relative diff 0.0). Fixed afterwards with `zlib.crc32` (regression test added). Results were **not** regenerated. *Effect:* none on reported numbers; re-runs will differ numerically

Later stages (replication, controls, tool evaluation)

V1 Exp B shift sign convention (bug found by independent audit (F1), fixed **before Exp B was analysed**). `shift = y6000 − y1500` is negative when NL improves; B2 "shift toward NL < 0.5" would have passed trivially. Replaced by `shift_toward_NL = y1500 − y6000`; unit test added. *Effect:* none (B not yet analysed)

V2 A2 baseline p0 miscounted (factual error in addendum v2.1 found by independent audit, verified). p0 stated 2/20; correct 3/20 (burgers_L1 f1 $= -1.9972$, $\bar{y} = 2.25$). Amended in addendum v2.2 **before A2 ran**: p0 $= 0.15$, success k ≥ 4 of 7. *Effect:* none (A2 not yet run)

V3 Diagnostic verdict wording (re-characterisation). "Not replicated" $\rightarrow$ *pre-registered criteria not met; inconclusive; carried by one family (nlheat 88.3 % of P2 gain; $\rho = +0.49$ without it); confounded with p (Spearman 0.708). Effect:* none (numbers unchanged)

V4 Stage chronology (disclosure). Stage 3 (addendum v2.1) was committed within minutes after the Exp A and Exp D results; Stage 4 (addendum v2.2) after two independent audits. All preregistrations are on a self-hosted git history; git tags could not be pushed to the remote. *Effect:* none

V5 Shared search draws (design property found by audit (F2/F3)). The random-search RNG depends only on (seed, method): the same 7 sampled configs are used for every instance and N within a seed (3 distinct sets per method in the whole study). v1 additionally used a process-salted hash (D5). *Effect:* none on verdicts; changes how "tuning" must be described; sensitivity analysis planned

V6 Neural batch/step count (design property found by audit (F02)). batch = N for N ≤ 512 (1500 steps) but 512 for N $= 1024$ (3000 steps): the neural step count doubles at N $= 1024$. Exp F removes this (full batch). *Effect:* none on verdicts

V7 Test/doc fixes (housekeeping). stale test count fixed (36 instances); weaker leakage tests replaced; BLAS thread env set at the top of entry scripts. *Effect:* none

V9 SurroSelect smoke-test bug (before any registered Stage-6 run) — 2026-11-02. During a smoke test of the tool on the new problem `lshape_L1` (seed 0, N = 64 and 256) and a quick check of the KRR and polynomial members on `heatcomp_L2` and `convdiff_L1` (N = 64, 256), the wide-range kernel-ridge member returned a test error of 0.5 (10x worse than any other member). Cause: the closed-form leave-one-out residual divides by 1 - H_ii, which was clipped at 1e-6 (and computed by cancellation), producing spuriously tiny LOO scores for near-interpolating fits; the length-scale grid was also defined for inputs on [0,1] but with the [-1,1] limits of the benchmark's wide-range control. Fix (committed before any registered run): 1 - H_ii computed without cancellation (KRR: sum_j Q_ij^2 lam N/(w_j + lam N); polynomial: (1 - sum U_ij^2) + sum U_ij^2 lam N/(s_j^2 + lam N)), no clipping above 1e-12, ridge values with condition number above 1e12 skipped, length-scale grid [0.05, 15]·sqrt(p) on [0,1]-scaled inputs (= [0.1, 30]·sqrt(p) on [-1,1]); a regression test was added. The smoke tests showed test errors of all pool members on lshape_L1 (seed 0, N = 64, 256) and of the KRR and polynomial members on three problems; the pool, criteria, tie band, folds and all pre-registered thresholds were NOT changed after seeing them. Addendum v3.0 stated that no surrogate error on the new problems had been seen; this entry corrects that statement.

V10 SurroSelect pickling bug found by the plain-file validation (2026-11-04). After the `metric` option was added (v3.1), a fitted selector stored a lambda and could no longer be saved by `surroselect fit --model`; the command-line unit test had been run before the change and the full tool test suite was not re-run. The plain-file validation (experiment E8) exposed the failure ("Can't get local object"). Fix: the metric function is a property computed from the metric name. The registered evaluation (experiments C, S, R) uses the library interface without pickling and is unaffected; all tool tests were re-run after the fix.

V11 Retrospective run promised in addendum v3.0 section 6.4 was reported late (2026-11-05). Review round 1 noted that the descriptive run of the tool on the 36 earlier problems (N = 256, seed 0), promised in addendum v3.0 section 6.4, had not been delivered. It was made on 2026-

11-05 (`code/experiments/exp_retro.py`, `analyze_retro.py`) and is reported in Supplement section S-retro; it is design-informed and not evidence for criteria S1-B2. No pool, grid or criterion was changed.

V12 Tuned-network search space narrowed in the addendum text after the criteria commit (2026-11-02). Commit 781e318, ten minutes after a99c2c6, changed the addendum's description of the tuned network from depth 2-5, width 64-512 to depth 2-4, width 64-256 and added the 6,000-epoch final training, for compute reasons, before any evaluation run (chain script 08:51). This is disclosed here and in the paper's limitations; it weakens the tuned network.

V13 Scalar-output prediction shape bug in the tool (found by review 03, 2026-11-05). For outputs with a single column, `GPFull.predict` and `PodGP.predict` returned 1-D arrays that broadcast against (n, 1) targets, giving meaningless NRMSE for full-field GP members on the real datasets (20-74). Fixed (predictions reshaped to (n, m)), regression test added. The real-data runs were re-checked with the fixed tool (`results/s6/real_recheck_after_shape_fix.json`; see the paper's real-data paragraph).

V14 Boundary-region bug in one new problem found by unit tests (2026-11-05). Before any Stage 7 surrogate was fitted: the first version of `make_disc_elast` selected the clamped and loaded arcs with `facets_satisfying` on a coordinate condition of facet midpoints; that picks interior facets inside the wedge as well, so a whole wedge was clamped and loaded, and the solution did not converge under mesh refinement (relative change between refinement levels about 0.5). The unit test for mesh convergence exposed it. Fix: select boundary facets only, by polar angle of the facet midpoint, and place the probes at 41 fixed physical points (not at mesh nodes of one particular level); mesh level raised from 4 to 5. A regression test asserts that the centre probe moves. Nothing had been fitted to this problem. Effect: none on earlier results (the problem is new).

V15 Documentation corrections and post hoc analyses after review 04 (2026-11-06). No run, criterion or result was changed. (a) Addendum v3.2 §7.4 says that the comparison with the benchmark's own selection (S3) was "only just met" in Stage 6; it was **not** met there (median ratio 1.01; 11 better, 9 worse). The sentence is a documentation error in a prediction, not a criterion; the paper's prediction table says so. (b) The one-line labels of the two Helmholtz problems in the problem

module ("near resonance", "several resonances") are inaccurate: the range of k^2 reaches the first eigenvalue $2\pi^2 \approx 19.7$ of the unit square at its upper end (L1) and contains the double eigenvalue $5\pi^2 \approx 49.3$ (L2); the paper's table uses corrected text, the code labels were left unchanged because the code is frozen. (c) the aggregation script for Stage 7 (which formats macros and tables) was changed after the freeze: the pooled column of the criteria table carries numbers only, because the thresholds were set for eight problems; the count of met verdicts was corrected (five of seven evaluable). (d) a post hoc analysis script was added after the results and is labelled post hoc in the paper: a carry-forward baseline for the reach-target classification, H1a/H1c without the two Helmholtz problems, per-job worst regret, error levels on the Helmholtz problems, the tool time without the evaluation-only the refit of all members, and the best fixed member in E7. (e) The unit tests of the new solvers were strengthened (positivity, resonance, sector sensitivity). No pool member, grid, threshold or data protocol changed.

S6. Data-leakage audit

Training, test and extrapolation inputs are drawn from separate named random streams. Over all 44 problems the number of exact input matches between the training pool and the test or extrapolation sets is 0. A planted-leak test (20 test samples copied into the training set of problem `heatcomp_L1`, $N = 64$) is flagged by the overlap check (20 exact matches, none without the leak), and the error on the leaked samples is 229 times lower than on the other test samples, showing that a leak would be visible. Validation errors do not depend on test labels (unit test, corrupting all test labels). Fitting steps (POD, scalers, hyper-parameters, cross-validation folds) use training runs only. Two properties remain: training subsets are nested within a seed, and seeds share the pool and test set, so seed variance understates sampling variance at large N. Comparisons in the mechanism study that use the best network by *test* error are oracle comparisons and are labelled as such.

S7. Per-problem results

Table S5 gives $\bar{y}$ (reduced basis against direct networks, validation-selected members, seed means) for every problem; positive values favour the reduced basis.

Table S5: $\bar{y}$ per problem and budget (positive: reduced basis better).

Problem	p	N=64	N=256	N=1024
advdiff_L1	4	-0.04	0.17	0.21
advdiff_L2	4	-0.09	-0.00	-0.07
bratu_L1	3	2.78	2.98	3.31
bratu_L2	3	2.14	2.97	3.08
burgers_L1	4	-0.06	0.68	2.25
burgers_L2	4	0.01	0.08	0.11
convdiff_L1	3	-0.03	0.60	0.48
convdiff_L2	4	-0.05	0.10	0.04
darcy_L1	4	0.15	1.12	1.13
darcy_L2	10	-0.00	-0.05	-0.19
darcyhd_L1	16	0.04	-0.02	-0.13
darcyhd_L2	32	0.03	0.00	-0.07
darcylin_L1	16	0.18	0.63	0.48
darcylin_L2	32	0.04	0.43	0.18
diff1d_L1	4	0.89	2.36	2.70
diff1d_L2	8	0.17	0.62	1.00
fisher_L1	4	-0.21	0.09	0.52
fisher_L2	4	-0.26	-0.26	-0.44
heat2d_L1	4	0.18	0.93	1.38
heat2d_L2	4	-0.04	-0.10	-0.34
heatcomp_L1	4	0.00	0.52	0.88
heatcomp_L2	9	-0.01	-0.06	-0.12
heatmove_L1	3	0.13	0.13	-0.17
heatmove_L2	5	-0.04	-0.02	-0.15
heatmv3_L1	3	0.12	-0.11	-0.44
heatp_L1	3	0.77	2.49	2.87
heatp_L2	3	1.08	2.71	2.90
helm_L1	3	1.20	2.69	2.93
helm_L2	3	-0.16	2.08	2.64
incl3_L1	3	0.02	0.05	-0.21
incl_L1	4	-0.02	-0.00	-0.25
incl_L2	4	0.02	-0.01	-0.13
kdv_L1	4	-0.14	-0.26	-0.23
kdv_L2	4	-0.04	-0.03	0.29
lshape_L1	4	0.20	1.14	1.66
lshape_L2	7	-0.11	0.01	-0.01
nlheat_L1	3	1.25	2.51	2.72
nlheat_L2	3	0.75	2.16	2.58
reacdiff_L1	3	0.85	1.62	2.57
reacdiff_L2	3	0.57	1.96	2.38
transport_L1	4	-0.26	0.05	0.75
transport_L2	4	-0.32	-0.23	-0.33
wave_L1	4	0.16	0.29	0.78
wave_L2	4	0.10	-0.06	-0.04

S8. Descriptive run on the earlier problems

Addendum v3.0 (Sec. 6.4) promised a descriptive run of the tool on the earlier problems, with one seed and $N = 256$. The run was not reported in the first version of this work and was made afterwards (this is deviation V11 in the log). It is *design-informed*: the pool and the grids of the tool were chosen with knowledge of these problems, so the run is not evidence for the criteria. On the 36 earlier problems solved with our own code, the median regret against the oracle (best pool member by test error) was 1.00 (largest 9.25; 97% of the problems within a factor two); the largest regrets were on `fisher_L1` (9.2), `advdiff_L1` (1.7), `incl3_L1` (1.2). The tool's error was a median 0.99 times that of the benchmark's validation-selected method (better by more than 10% on 16 problems, worse on 7), and 0.88 times that of the benchmark's POD+GP. The basis flag was raised on 4 of the problems.

S9. The tool

Installation and use. The package `surroselect` (directory `surroselect/` of the repository; MIT licence, copyright holder to be set by the authors) depends on NumPy, SciPy and scikit-learn; PyTorch and Optuna are needed only for the optional `TunedMLP`, which the command line adds to the pool with `-with-network`. In Python:

```
sel = SurrogateSelector(cv=4).fit(X, Y)  # X: runs x inputs
sel.learning_curve()
print(make_report(sel, target=0.01))
Y_new = sel.predict(X_new)
```

On the command line, with plain `.npz` files that hold arrays `X` and `Y`:

```
surroselect fit train.npz --model m.pkl --target 0.01 --report report.md
surroselect predict m.pkl X_new.npy --out Y_pred.npy
surroselect evaluate m.pkl test.npz
```

Default pool. POD+GP (Matérn-5/2, rank 32), POD+GP (squared exponential, rank 32), POD+GP (Matérn-5/2, rank 8), POD+KRR with a wide grid of kernel, length-scale and ridge selected by closed-form leave-one-out error, POD+polynomial (degree up to six, at most 3000 features, leave-one-out), and a GP on the full field (Matérn-5/2 or squared exponential; only for $m \leq 2000$). An optional metric, the normalised RMSE $\|Y - \hat{Y}\|_F / \|Y - \bar{Y}_{\text{train}}\|_F$, is available for outputs with a large offset.

Validation of the tool on plain files. To check the tool itself, without the benchmark pipeline in the loop, three of the new problems (`heatcomp_L1`, `lshape_L1`, `convdiff_L1`; $N = 256$) were exported to `.npz` files and processed with the three command-line commands as separate processes. All 3 of 3 sequences ran without error, and the error printed by `evaluate` agreed with an independent computation from the saved predictions to within 10^{-12} in 3 of 3 cases (`results/s6/plain_files_validation.json`).

Tests. 11 unit tests for the tool (closed-form leave-one-out against brute force and its stability near interpolation, exact polynomial recovery, selection accuracy, the basis-limit flag, the power-law inversion, the report, the metric option and the command-line round trip), 6 tests for the finite-element problems and the 67 tests of the benchmark package are collected by the repository; `results/logs/tests_*.log` records the final runs. A fitted selector could not be saved in an early version after the metric option was added (deviation V10); the plain-file validation found it.

Known limits. The tool assumes exchangeable runs (A3); it does not model noise; for $p \gtrsim 10$ the default GP kernels have many hyper-parameters fitted on at most 300 runs; the forecast is an extrapolation without guarantee; and the optional neural member uses a small search budget.